




Efficient and scalable graph generation through iterative local expansion

Luka Lafaye de Micheaux

August 21, 2025

Introduction

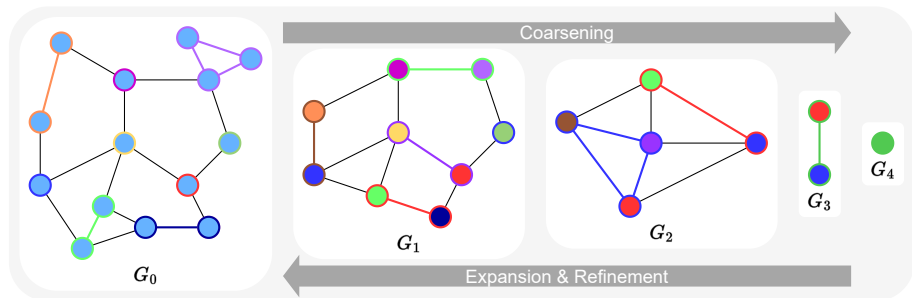


Source: Bachechi et al. 2022 (Road Network Graph), Susi et al. 2021 (FNS Neuron model), Wikipedia Commons, Tao et al. 2025 (CGM model)

Problem statement

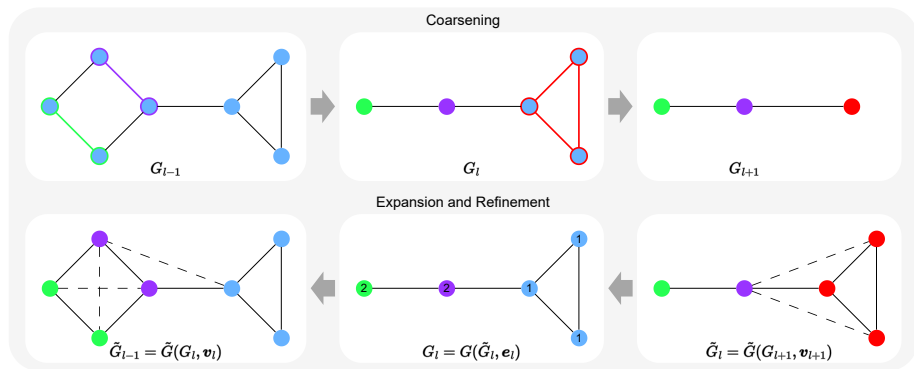
- Given real-world graphs sampled from an unknown distribution $\mathcal{G}_{\text{real}}$, we want to learn a generative model $p_{\theta}(G)$ that approximates this distribution.
- The sampled new graphs $G \sim p_{\theta}$ from the model should be similar to those from $\mathcal{G}_{\text{real}}$.
- Challenges:
 - Modeling the full joint distribution over all node pairs,
 - Capturing both local patterns and global structure,
 - Scaling to large graphs and generalizing to unseen sizes.

Iterative coarsening, expansion, and refinement



Source: Bergmeister et al. 2024 - Illustration of the iterative coarsening, expansion, and refinement process.

Notations of intermediate states



Source: Bergmeister et al. 2024 - Notations of intermediate steps during the process.

Likelihood of a graph

- Coarsening provides a sequence $G^{(0)} = G, G^{(1)}, \dots, G^{(L-1)}, G^{(L)}$ where $G^{(L)}$ is a single node graph
- Expansion and refinement provides the inverse sequence $G^{(L)}, \tilde{G}^{(L-1)}, G^{(L-1)}, \dots, \tilde{G}^{(1)}, G^{(1)}, \tilde{G}^{(0)}, G^{(0)} = G$
- Let $\omega \in E(G)$ denote an expansion-refinement sequence that reconstructs G , and $\pi \in C(G)$ a coarsening sequence ending in a singleton.
- Then the likelihood of G is defined as:

$$p(G) = \sum_{\omega \in E(G)} p(\omega) > \sum_{\pi \in C(G)} p(\pi)$$

Contraction families and elbow trick

- We restrict possible contraction sets to belong to $\mathcal{F}(G)$ (edge contractions in my implementation), and coarsening sequences are sampled in $\mathcal{C}_{\mathcal{F}}(G)$. Hence:

$$p(G) = \sum_{\omega \in E(G)} p(\omega) > \sum_{\pi \in \mathcal{C}(G)} p(\pi) > \sum_{\pi \in \mathcal{C}_{\mathcal{F}}(G)} p(\pi)$$

- Given a distribution $q(\pi \mid G)$ over coarsening sequences $\mathcal{C}_{\mathcal{F}}(G)$,

$$p(G) > \mathbb{E}_{\pi \sim q(\pi \mid G)} \left[\frac{p(\pi)}{q(\pi \mid G)} \right]$$


- Using the elbow trick we have:

$$\log(p(G)) > \mathbb{E}_{\pi \sim q(\pi \mid G)} [\log(p(\pi)) - \log(q(\pi \mid G))]$$


Implementation details

- Sampling $\pi \sim q(\pi \mid G)$ provides coarsening sequences from which we can extract $G^{(L)}, \tilde{G}^{(L-1)}(v_L), G^{(L-1)}(e_{L-1}), \dots, \tilde{G}^{(1)}(v_2), G^{(1)}(e_1), \tilde{G}^{(0)}(v_1), G^{(0)}(e_0) = G$
- The dataset contains pairs of (v_l, e_l) stored in $\tilde{G}^{(l)}$, resp. as node and edge features. They are noised into (\hat{e}_l, \hat{v}_l) during training.
- The denoising model learns to recover (e_l, v_l) given $\tilde{G}^{(l)}$. It consists of two GCN layers followed by two MLP heads, one used to recover v_l , and the other used to recover e_l .
- Implementation was done from scratch

Simplified edge selection cost function



Before merge



After merge

- Each eigenvector of the Laplacian corresponds to a mode
- Signals x get propagated in the graph using $D^{-1}A \times x$
- Each signal can be decomposed in modes (one per eigenvector). $e^l(i)$ corresponds to how much node i influences mode of eigenvector e^l
- Let all eigenvectors of L be noted e^l . We use:

$$\text{Cost}(i, j) = \frac{1}{N-1} \sum_{l=1}^{N-1} (e^l(i) - e^l(j))^2.$$

Basic model architecture

- Three GCN layers

$$h^{(1)} = \text{ReLU}(\text{GCNConv}_1(\hat{v}_l)),$$

$$h^{(2)} = \text{ReLU}(\text{GCNConv}_2(h^{(1)})),$$

$$h^{(3)} = \text{ReLU}(\text{GCNConv}_3(h^{(2)})).$$

- Node noise head

$$\hat{\epsilon}_v = \text{MLP}_{\text{node}}(h^{(3)})$$

- Edge noise head

$$\hat{\epsilon}_e = \text{MLP}_{\text{edge}}([h_i^{(3)}, h_j^{(3)}, e_{ij}^{(t)}])$$


- GCN Reminder: $H^{(l+1)} = \sigma(\tilde{D}^{-\frac{1}{2}} \tilde{A} \tilde{D}^{-\frac{1}{2}} H^{(l)} W^{(l)})$ where $\tilde{A} = A + I$, $\tilde{D}_{ii} = \sum_j \tilde{A}_{ij}$.

Improving model w.r.t. L2 loss on Erdos-Renyi dataset


- L2 loss on random noise sampled $X \sim \mathcal{N}(0, 1)$,
 $P(|X| \leq 3) = 99.7\%$.
- Initial learning rate $10e-5$ and 60 epochs obtained average loss of 0.97.
- Adjusted learning rate to $10e-4$ and obtained average loss 0.85
- Added sinusoidal encodings instead of linearly projecting timestep scalar and got average loss of 0.70
- Added norm layers after each GCN conv and dropout after last conv and obtained average loss of 0.65

- Erdos-Renyi synthetic dataset ($n=10$, $p=0.5$)
 - Average degree
 - Average clustering coefficients
 - Average shortest path
 - Degree distribution
- Planar graphs dataset using plantri ($c=3m$ $n=12$)
 - Density
 - Planarity proportion

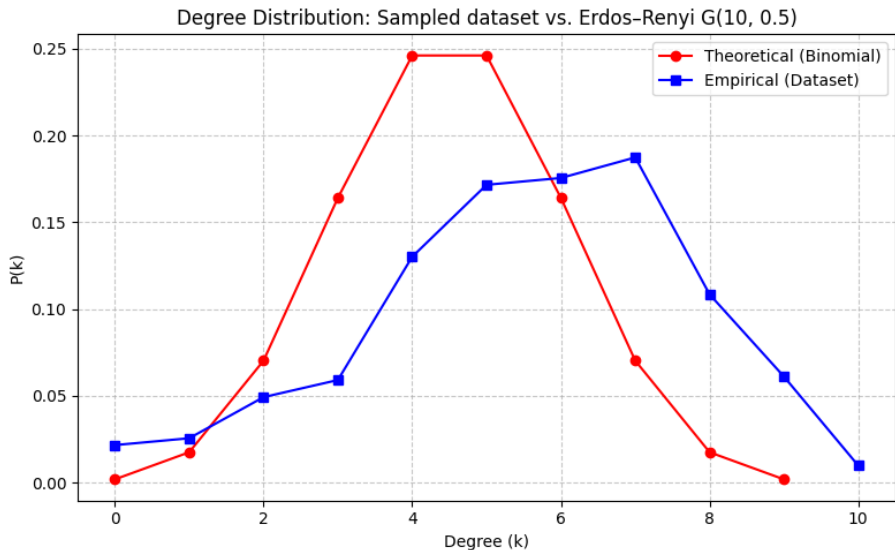
Training loss on Erdos-Renyi dataset



Training loss on planar dataset



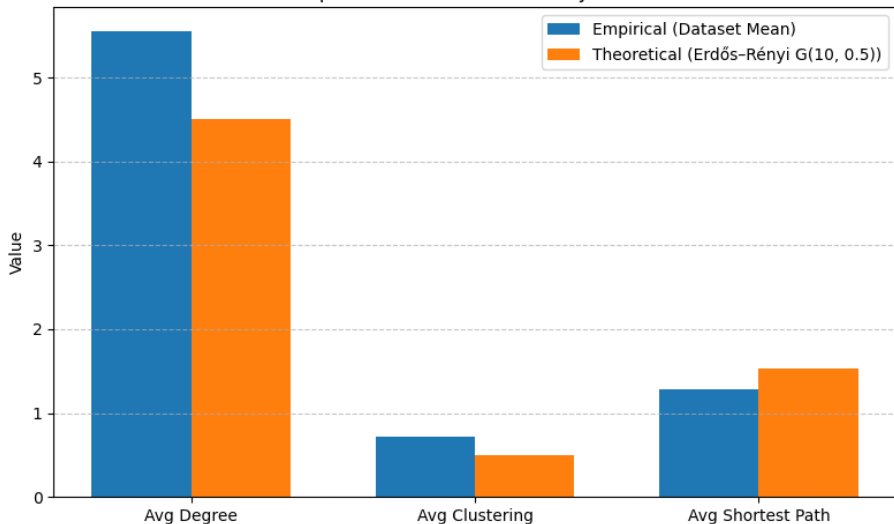
Results - Erdos Renyi



Problem with self edges

Results - Erdos Renyi

Sampled dataset vs. Erdos-Renyi metrics



- Proportion of planar graphs: 10%
- Average density: 0.410 (num edges 41% of complete graph)
- Harder task, bad results explained by lack of edge features

Bonus: improving results

- Improve edge selection cost function, use neighborhoods coarsening instead of edge coarsening
- Try Edge-GAT and leverage edge features
- Randomly skip/connect nodes during message passing
- Denoise graphs while training and add loss term that penalizes crossings?